

# Towards Group-Conditional Coverage Guarantees via Group-Invariant Representations

**Paul Melki**

*EXXACT Robotics, Villeurbanne, France*

PAUL.MELKI@EXXACT-ROBOTICS.COM

**Lionel Bombrun**

**Jean-Pierre Da Costa**

*IMS Laboratory, Bordeaux Sciences Agro, University of Bordeaux, France*

LIONEL.BOMBRUN@IMS-BORDEAUX.FR

JEAN-PIERRE.DACOSTA@IMS-BORDEAUX.FR

**Editors:** Ernst Ahlberg, Ulf Johansson, Henrik Boström, Alberto Carlevaro, Johan Hallberg Szabadváry and Lars Carlsson

## Abstract

Split-conformal prediction is a solid framework that transforms any point predictor into a conformal set predictor with marginal coverage guarantees. In cases where observations are divided into different groups, these “average” guarantees may not be enough. Group-conditional coverage cannot be maintained, by default, using marginal calibration since different groups have different distributions of nonconformity scores. In this short paper, we propose a methodology for obtaining equalized group-conditional coverage via marginal calibration by learning group-invariant representations that disentangle the relationship between group-specific information and the nonconformity scores.

**Keywords:** Mondrian conformal prediction, group invariance, metadata normalization

## 1. Introduction

Suppose that we have a deep neural network trained for a classification task  $f : \mathcal{X} \rightarrow \mathcal{Y} = \{1, \dots, K\}$ ,  $K \in \mathbb{N}^*$ , in a supervised way on a finite set of labelled observations  $\{(\mathbf{x}_i, y_i), i = 1, \dots, n\}$ . Let  $\mathfrak{X}$  be the representation of the data learned by the neural network, and  $p^y$  be the estimated probability of class  $y$  (for example, the calibrated softmax value of  $y$ ). The split-conformal procedure (Papadopoulos et al., 2002; Vovk et al., 2022) is a powerful method to “transform” the class predictor  $f$  into a set predictor  $\mathcal{C}_\alpha : \mathcal{X} \rightarrow 2^{\mathcal{Y}}$  where  $2^{\mathcal{Y}}$  is the power set of  $\mathcal{Y}$  with the valid *marginal* coverage guarantee:

$$\mathbb{P}(Y_{n+1} \in \mathcal{C}_\alpha(\mathbf{X}_{n+1})) \geq 1 - \alpha, \quad (1)$$

where  $\alpha \in (0, 1)$  is a chosen tolerance to error. To do this, a non-conformity score function  $\Delta : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$  that quantifies the *strangeness* of any arbitrary couple  $(\mathbf{x}, y)$  needs to be defined, often based on the estimated probabilities  $p^y$  output by the neural network (Johansson et al., 2017; Melki et al., 2024).


The conformal calibration step consists of estimating  $\hat{q}$  as the  $\frac{\lceil (m+1)(1-\alpha) \rceil}{m}$  quantile of the scores computed on a held-out calibration set of size  $m$ , and the prediction step creates, for a new example  $\mathbf{x}_{n+1}$ , the prediction set of classes  $\mathcal{C}_\alpha(\mathbf{x}_{n+1}) = \{y : \Delta(\mathbf{x}_{n+1}, y) \leq \hat{q}\}$ .

## 2. Group-Conditional Conformal Prediction

In practice, we may have at our disposal, in addition to the labeled observations  $(\mathbf{x}, y)$ , some auxiliary data  $\mathbf{a} \in \mathcal{A}$  that provide additional information about the examples ( $\mathcal{A}$  is considered discrete and finite). For instance,  $\mathbf{a}$  may be a vector of variables describing the environmental conditions under which the example  $\mathbf{x}$  was observed. We would thus like to provide the following *group-conditional* coverage guarantee:

$$\mathbb{P}(Y_{n+1} \in \mathcal{C}_\alpha(\mathbf{X}_{n+1}) \mid \mathbf{a}_{n+1} = a) \geq 1 - \alpha, \forall a \in \mathcal{A}. \quad (2)$$

The split-conformal prediction procedure described in the introduction is not enough to provide the group-conditional guarantee in Equation (2). This is mainly due to the influence of exogenous factors, some of which may be encoded in  $\mathbf{a}$ , that result in each group having a different  $1 - \alpha$  quantile,  $\hat{q}_{\mathbf{a}_1} \neq \dots \neq \hat{q}_{\mathbf{a}_t} \neq \hat{q}$  where  $t = |\mathcal{A}|$  is the number of groups (panel (a) in the adjacent figure), thus requiring group-conditional calibration. Therefore, the question we ask is: Can we obtain equalized conditional quantiles so that marginal calibration is enough to maintain conditional coverage?




## 3. Group-Invariant Representations

To answer the previous question, we propose an intervention that consists of breaking the chain of influence of  $\mathbf{A}$  on the rest of the calibration procedure at the level of model training by making the model learn representations  $\mathbf{x}^\perp$  that are *invariant to the auxiliary information*  $\mathbf{A}$ , instead of  $\mathbf{x}$ , as shown in panel (b) in the figure above.

What we wish to achieve is the invariance of the nonconformity scores (that we will simply denote as  $\Delta$ ) with respect to the auxiliary information, conditional on the learned representations. In other words, we want to learn  $\mathbf{x}^\perp$  such that  $\Delta \perp \mathbf{A} \mid \mathbf{x}^\perp$ . If that is achieved, then we would have  $P(\Delta \mid \mathbf{x}^\perp, \mathbf{A}) = P(\Delta \mid \mathbf{x}^\perp)$ .

Metadata Normalization (MDN), introduced by Lu et al. (2021), is a method that aims at learning model representations that are orthogonal to some auxiliary data. MDN is a simple layer that can be integrated into any neural network, at different stages, in order to “filter out” the impact of exogenous variables.

Consider the figure to the right showing a conceptual representation of a MDN layer inserted between two consecutive layers. The MDN layer essentially tries to model the representation  $\mathbf{x}$  as a linear function of the auxiliary data by learning the linear model  $\mathbf{x} = \mathbf{A}\beta + \mathbf{r}$ , where  $\mathbf{r}$  is the residual term and  $\beta$  are linear parameters to be estimated. This model tries to separate  $\mathbf{x}$  into two parts:  $\mathbf{A}\beta$ , the component of  $\mathbf{x}$  that is linearly related to the metadata  $\mathbf{A}$ , and the residual term  $\mathbf{r} = \mathbf{x}^\perp$ , which is orthogonal to the linear subspace spanned by the metadata. By integrating multiple MDN layers in the network’s architecture we would effectively disentangle the nonconformity scores from the auxiliary data. Preliminary results on the agricultural task of crop and weed image classification with environmental auxiliary data and MDN integration will be shown during the poster session.



## References

- Ulf Johansson, Henrik Linusson, Tuve Löfström, and Henrik Boström. Model-agnostic nonconformity functions for conformal classification. In *International Joint Conference on Neural Networks (IJCNN)*, page 2072–2079, May 2017. doi: 10.1109/IJCNN.2017.7966105.
- Mandy Lu, Qingyu Zhao, Jiequan Zhang, Kilian M. Pohl, Li Fei-Fei, Juan Carlos Niebles, and Ehsan Adeli. Metadata normalization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 10917–10927, June 2021.
- Paul Melki, Lionel Bombrun, Boubacar Diallo, Jérôme Dias, and Jean-Pierre Da Costa. The penalized inverse probability measure for conformal classification. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops*, Seattle WA USA, 2024.
- Harris Papadopoulos, Kostas Proedrou, Volodya Vovk, and Alex Gammerman. Inductive confidence machines for regression. In Tapio Elomaa, Heikki Mannila, and Hannu Toivonen, editors, *Machine Learning: ECML 2002*, Lecture Notes in Computer Science, page 345–356, Berlin, Heidelberg, 2002. Springer. ISBN 978-3-540-36755-0. doi: 10.1007/3-540-36755-1\_29.
- Vladimir Vovk, Alexander Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer International Publishing, Cham, 2022. ISBN 978-3-031-06648-1. doi: 10.1007/978-3-031-06649-8.